

line_word_segmentation

September 18, 2021

0.0.1 Segmentacja słów i lini

Ten notatnik demonstruje algorytm pozwalający na zapewnienia ramek ograniczających otaczających odręcznie zapisany tekst.

Dane wejściowe: Obraz zawierający odręczny tekst.

Dane wyjściowe: lista ramek ograniczających.

```
[2]: import multiprocessing
import os
import random
import time

import cv2
import matplotlib.pyplot as plt
import matplotlib.patches as patches

import mxnet as mx
from mxnet import nd, autograd, gluon
from mxnet.image import resize_short
from mxboard import SummaryWriter
from mxnet.gluon.model_zoo.vision import resnet34_v1
from mxnet.contrib.ndarray import MultiBoxPrior, MultiBoxTarget,
    MultiBoxDetection, box_nms
import numpy as np
from skimage.draw import line_aa
from skimage import transform as skimage_tf

np.seterr(all='raise')

mx.random.seed(42)

from ocr.utils.iam_dataset import IAMDataset
from ocr.utils.draw_box_on_image import draw_boxes_on_image
```

Definiowanie sieci Do identyfikacji każdej linii używamy sieci SSD.

```
[4]: class SSD(gluon.Block):
    def __init__(self, num_classes, ctx, **kwargs):
        super(SSD, self).__init__(**kwargs)

        # Seven sets of anchor boxes are defined. For each set, n=2 sizes and
        ↪ m=3 ratios are defined.
        # Four anchor boxes (n + m - 1) are generated: 2 square anchor boxes
        ↪ based on the n=2 sizes and 2 rectangles based on
        # the sizes and the ratios. See https://discuss.mxnet.io/t/question-regarding-ssd-algorithm/1307 for more information.

        #self.anchor_sizes = [[.1, .2], [.2, .3], [.2, .4], [.4, .6], [.5, .7],
        ↪ [.6, .8], [.7, .9]]
        #self.anchor_ratios = [[1, 3, 5], [1, 3, 5], [1, 6, 8], [1, 5, 7], [1,
        ↪ 6, 8], [1, 7, 9], [1, 7, 10]]

        self.anchor_sizes = [[.1, .2], [.2, .3], [.2, .4], [.3, .4], [.3, .5],
        ↪ [.4, .6]]
        self.anchor_ratios = [[1, 3, 5], [1, 3, 5], [1, 6, 8], [1, 4, 7], [1,
        ↪ 6, 8], [1, 5, 7]]

        self.num_anchors = len(self.anchor_sizes)
        self.num_classes = num_classes
        self.ctx = ctx
        with self.name_scope():
            self.body, self.downsamples, self.class_preds, self.box_preds =
            ↪ self.get_ssd_model()
            self.downsamples.initialize(mx.init.Normal(), ctx=self.ctx)
            self.class_preds.initialize(mx.init.Normal(), ctx=self.ctx)
            self.box_preds.initialize(mx.init.Normal(), ctx=self.ctx)

    def get_body(self):
        """
        Create the feature extraction network of the SSD based on resnet34.
        The first layer of the res-net is converted into grayscale by averaging
        ↪ the weights of the 3 channels
        of the original resnet.

        Returns
        -----
        network: gluon.nn.HybridSequential
            The body network for feature extraction based on resnet

        """
        pretrained = resnet34_v1(pretrained=True, ctx=self.ctx)
        pretrained_2 = resnet34_v1(pretrained=True, ctx=mx.cpu(0))
```

```

        first_weights = pretrained_2.features[0].weight.data().mean(axis=1).
→expand_dims(axis=1)
        # First weights could be replaced with individual channels.

        body = gluon.nn.HybridSequential()
        with body.name_scope():
            first_layer = gluon.nn.Conv2D(channels=64, kernel_size=(7, 7),
→padding=(3, 3), strides=(2, 2), in_channels=1, use_bias=False)
            first_layer.initialize(mx.init.Normal(), ctx=self.ctx)
            first_layer.weight.set_data(first_weights)
            body.add(first_layer)
            body.add(*pretrained.features[1:-3])
        return body

    def get_class_predictor(self, num_anchors_predicted):
        """
        Creates the category prediction network (takes input from each
→downsampled feature)

        Parameters
        -----

        num_anchors_predicted: int
            Given n sizes and m ratios, the number of boxes predicted is n+m-1.
            e.g., sizes=[.1, .2], ratios=[1, 3, 5] the number of anchors
→predicted is 4.

        Returns
        -----

        network: gluon.nn.HybridSequential
            The class predictor network
        """
        return gluon.nn.Conv2D(num_anchors_predicted*(self.num_classes + 1),
→kernel_size=3, padding=1)

    def get_box_predictor(self, num_anchors_predicted):
        """
        Creates the bounding box prediction network (takes input from each
→downsampled feature)

        Parameters
        -----

        num_anchors_predicted: int
            Given n sizes and m ratios, the number of boxes predicted is n+m-1.

```

e.g., sizes=[.1, .2], ratios=[1, 3, 5] the number of anchors
 ↳ predicted is 4.

Returns

pred: gluon.nn.HybridSequential

The box predictor network

'''

pred = gluon.nn.HybridSequential()

with pred.name_scope():

↳ pred.add(gluon.nn.Conv2D(channels=num_anchors_predicted*4,
 ↳ kernel_size=3, padding=1))

return pred

def get_down_sampler(self, num_filters):

'''

*Creates a two-stacked Conv-BatchNorm-Relu and then a pooling layer to
 downsample the image features by half.*

'''

out = gluon.nn.HybridSequential()

for _ in range(2):

out.add(gluon.nn.Conv2D(num_filters, 3, strides=1, padding=1))

out.add(gluon.nn.BatchNorm(in_channels=num_filters))

out.add(gluon.nn.Activation('relu'))

out.add(gluon.nn.MaxPool2D(2))

out.hybridize()

return out

def get_ssd_model(self):

'''

*Creates the SSD model that includes the image feature, downsample,
 ↳ category*

and bounding boxes prediction networks.

'''

body = self.get_body()

downsamples = gluon.nn.HybridSequential()

class_preds = gluon.nn.HybridSequential()

box_preds = gluon.nn.HybridSequential()

downsamples.add(self.get_down_sampler(32))

downsamples.add(self.get_down_sampler(32))

downsamples.add(self.get_down_sampler(32))

for scale in range(self.num_anchors):

↳ num_anchors_predicted = len(self.anchor_sizes[0]) + len(self.
 ↳ anchor_ratios[0]) - 1

```

        class_preds.add(self.get_class_predictor(num_anchors_predicted))
        box_preds.add(self.get_box_predictor(num_anchors_predicted))

    return body, downsamples, class_preds, box_preds

def ssd_forward(self, x):
    """
    Helper function of the forward pass of the ssd
    """
    x = self.body(x)

    default_anchors = []
    predicted_boxes = []
    predicted_classes = []

    for i in range(self.num_anchors):
        default_anchors.append(MultiBoxPrior(x, sizes=self.anchor_sizes[i],
        ratios=self.anchor_ratios[i]))
        predicted_boxes.append(self._flatten_prediction(self.
        box_preds[i](x)))
        predicted_classes.append(self._flatten_prediction(self.
        class_preds[i](x)))
        if i < len(self.downsamples):
            x = self.downsamples[i](x)
        elif i == 3:
            x = nd.Pooling(x, global_pool=True, pool_type='max', kernel=(4,
        4))

    return default_anchors, predicted_classes, predicted_boxes

def forward(self, x):
    default_anchors, predicted_classes, predicted_boxes = self.
    ssd_forward(x)
    # we want to concatenate anchors, class predictions, box predictions
    from different layers
    anchors = nd.concat(*default_anchors, dim=1)
    box_preds = nd.concat(*predicted_boxes, dim=1)
    class_preds = nd.concat(*predicted_classes, dim=1)
    class_preds = nd.reshape(class_preds, shape=(0, -1, self.num_classes +
    1))

    return anchors, class_preds, box_preds

def _flatten_prediction(self, pred):
    """
    Helper function to flatten the predicted bounding boxes and categories
    """
    return nd.flatten(nd.transpose(pred, axes=(0, 2, 3, 1)))

```

```

def training_targets(self, default_anchors, class_predicts, labels):
    """
    Helper function to obtain the bounding boxes from the anchors.
    """
    class_predicts = nd.transpose(class_predicts, axes=(0, 2, 1))
    box_target, box_mask, cls_target = MultiBoxTarget(default_anchors,
→labels, class_predicts)
    return box_target, box_mask, cls_target

```

Funkcja straty Sieć SSD został przeszkolona w celu zminimalizowania błędu klasyfikacji i wygładzonej utraty L1 między przewidywaną a rzeczywistą ramką ograniczającą. Gładka strata L1 jest zdefiniowana poniżej.

```

[6]: class SmoothL1Loss(gluon.loss.Loss):
    """
    A SmoothL1loss function defined in https://gluon.mxnet.io/
→chapter08_computer-vision/object-detection.html
    """
    def __init__(self, batch_axis=0, **kwargs):
        super(SmoothL1Loss, self).__init__(None, batch_axis, **kwargs)

    def hybrid_forward(self, F, output, label, mask):
        loss = F.smooth_l1((output - label) * mask, scalar=1.0)
        return F.mean(loss, self._batch_axis, exclude=True)

```

Transformacja i augmentacja danych Dane są przekształcane, aby obrazy i etykiety mogły zostać wprowadzone do sieci. Dane treningowe są powiększane, dzięki czemu obrazy są losowo tłumaczone, a linie są losowo usuwane.

```

[8]: def augment_transform(image, label):
    """
    1) Function that randomly translates the input image by +-width_range and
→+-height_range.
    The labels (bounding boxes) are also translated by the same amount.
    2) Each line can also be randomly removed for augmentation. Labels are also
→reduced to correspond to this
    data and label are converted into tensors by calling the "transform"
→function.
    """
    ty = random.uniform(-random_y_translation, random_y_translation)
    tx = random.uniform(-random_x_translation, random_x_translation)

    st = skimage_tf.SimilarityTransform(translation=(tx*image.shape[1],
→ty*image.shape[0]))
    image = skimage_tf.warp(image, st, cval=1.0)

```

```

    label[:, 0] = label[:, 0] - tx/2 #NOTE: Check why it has to be halved
    → (found experimentally)
    label[:, 1] = label[:, 1] - ty/2

    index = np.random.uniform(0, 1.0, size=label.shape[0]) > random_remove_box
    for i, should_output_bb in enumerate(index):
        if should_output_bb == False:
            (x, y, w, h) = label[i]
            (x1, y1, x2, y2) = (x, y, x + w, y + h)
            (x1, y1, x2, y2) = (x1 * image.shape[1], y1 * image.shape[0],
                                x2 * image.shape[1], y2 * image.shape[0])
            (x1, y1, x2, y2) = (int(x1), int(y1), int(x2), int(y2))
            x1 = 0 if x1 < 0 else x1
            y1 = 0 if y1 < 0 else y1
            x2 = 0 if x2 < 0 else x2
            y2 = 0 if y2 < 0 else y2
            image_h, image_w = image.shape
            x1 = image_w-1 if x1 >= image_w else x1
            y1 = image_h-1 if y1 >= image_h else y1
            x2 = image_w-1 if x2 >= image_w else x2
            y2 = image_h-1 if y2 >= image_h else y2
            image[y1:y2, x1:x2] = image[y1, x1]

    augmented_labels = label[index, :]
    return transform(image*255., augmented_labels)

def transform(image, label):
    """
    Function that converts resizes image into the input image tensor for a CNN.
    The labels (bounding boxes) are expanded, converted into (x, y, x+w, y+h),
    → and
    zero padded to the maximum number of labels. Finally, it is converted into
    → a float
    tensor.
    """
    max_label_n = 128 if detection_box == "word" else 13

    # Resize the image
    image = np.expand_dims(image, axis=2)
    image = mx.nd.array(image)
    image = resize_short(image, image_size)
    image = image.transpose([2, 0, 1])/255.

    # Expand the bounding box by expand_bb_scale
    bb = label.copy()
    new_w = (1 + expand_bb_scale) * bb[:, 2]

```

```

new_h = (1 + expand_bb_scale) * bb[:, 3]

bb[:, 0] = bb[:, 0] - (new_w - bb[:, 2])/2
bb[:, 1] = bb[:, 1] - (new_h - bb[:, 3])/2
bb[:, 2] = new_w
bb[:, 3] = new_h
label = bb

# Convert the predicted bounding box from (x, y, w, h to (x, y, x + w, y +
↪h)
label = label.astype(np.float32)
label[:, 2] = label[:, 0] + label[:, 2]
label[:, 3] = label[:, 1] + label[:, 3]

# Zero pad the data
label_n = label.shape[0]
label_padded = np.zeros(shape=(max_label_n, 5))
label_padded[:label_n, 1:] = label
label_padded[:label_n, 0] = np.ones(shape=(1, label_n))
label_padded = mx.nd.array(label_padded)
return image, label_padded

```

```

[11]: def run_epoch(e, network, dataloader, trainer, log_dir, print_name, is_train,
↪update_metric):
    """
    Run one epoch to train or test the SSD network

    Parameters
    -----

    e: int
        The epoch number

    network: nn.Gluon.HybridSequential
        The SSD network

    dataloader: gluon.data.DataLoader
        The train or testing dataloader that is wrapped around the iam_dataset

    log_dir: Str
        The directory to store the log files for mxboard

    print_name: Str
        Name to print for associating with the data. usually this will be
↪"train" and "test"

    is_train: bool

```


Boolean to indicate whether or not the CNN should be updated. is_train_
→ should only be set to true for the training data

Returns

network: gluon.nn.HybridSequential

The class predictor network

'''

```
total_losses = [0 for ctx_i in ctx]
for i, (X, Y) in enumerate(dataloader):
    X = gluon.utils.split_and_load(X, ctx)
    Y = gluon.utils.split_and_load(Y, ctx)

    with autograd.record(train_mode=is_train):
        losses = []
        for x, y in zip(X, Y):
            default_anchors, class_predictions, box_predictions = network(x)
            box_target, box_mask, cls_target = network.
            → training_targets(default_anchors, class_predictions, y)
            # losses
            loss_class = cls_loss(class_predictions, cls_target)
            loss_box = box_loss(box_predictions, box_target, box_mask)
            # sum all losses
            loss = loss_class + loss_box
            losses.append(loss)

        if is_train:
            for loss in losses:
                loss.backward()
            step_size = 0
            for x in X:
                step_size += x.shape[0]
            trainer.step(step_size)

        for index, loss in enumerate(losses):
            total_losses[index] += loss.mean().asscalar()

        if update_metric:
            cls_metric.update([cls_target], [nd.transpose(class_predictions,
            → (0, 2, 1))])
            box_metric.update([box_target], [box_predictions * box_mask])

            if i == 0 and e % send_image_every_n == 0 and e > 0:
                cls_probs = nd.SoftmaxActivation(nd.transpose(class_predictions,
            → (0, 2, 1)), mode='channel')
```

```

        output_image, number_of_bbs =
↳generate_output_image(box_predictions, default_anchors,
                                cls_probs,
↳box_target, box_mask,
                                cls_target, x,
↳y)
        print("Number of predicted {} BBs = {}".format(print_name,
↳number_of_bbs))
        with SummaryWriter(logdir=log_dir, verbose=False, flush_secs=5) as
↳sw:
            sw.add_image('bb_{}_image'.format(print_name), output_image,
↳global_step=e)

    total_loss = 0
    for loss in total_losses:
        total_loss += loss / (len(dataloader)*len(total_losses))

    with SummaryWriter(logdir=log_dir, verbose=False, flush_secs=5) as sw:
        if update_metric:
            name1, val1 = cls_metric.get()
            name2, val2 = box_metric.get()
            sw.add_scalar(name1, {"test": val1}, global_step=e)
            sw.add_scalar(name2, {"test": val2}, global_step=e)
            sw.add_scalar('loss', {print_name: total_loss}, global_step=e)

    return total_loss

```

Definowanie domyślny parametrów

```
[13]: detection_box = "word" # "word" or "line"
```

```
[14]: gpu_count = 4
expand_bb_scale = 0.05
min_c = 0.01
overlap_thres = 0.1 if detection_box == "line" else 0.001
topk = 150 if detection_box == "line" else 200

epochs = 20
learning_rate = 0.00005
batch_size = 32
image_size = 350

random_x_translation, random_y_translation = (0.03, 0.03) if detection_box ==
↳"line" else (0.005, 0.005)
random_remove_box = 0.1

```

```

log_dir = "./logs/line_word_segmentation"
checkpoint_dir, checkpoint_name = "model_checkpoint", "ssd_"+detection_box+"
    ↳params"

print_every_n = 5
send_image_every_n = 20
save_every_n = 50

```

Pętla treningowa

```

[23]: ctx = [mx.gpu(0)]

train_ds = IAMDataset("form_bb", output_data="bb",
    ↳output_parse_method=detection_box, train=True)
print("Number of training samples: {}".format(len(train_ds)))

test_ds = IAMDataset("form_bb", output_data="bb",
    ↳output_parse_method=detection_box, train=False)
print("Number of testing samples: {}".format(len(test_ds)))

train_data = gluon.data.DataLoader(train_ds.transform(augment_transform),
    ↳batch_size, shuffle=True, last_batch="rollover", num_workers=8)
test_data = gluon.data.DataLoader(test_ds.transform(transform), batch_size,
    ↳shuffle=False, last_batch="keep", num_workers=8)

```

```

Number of training samples: 967
Number of testing samples: 232

```

```

[24]: net = SSD(num_classes=2, ctx=ctx)
net.hybridize()

```

```

[25]: cls_loss = gluon.loss.SoftmaxCrossEntropyLoss()
box_loss = SmoothL1Loss()
cls_loss.hybridize()
box_loss.hybridize()

best_test_loss = 10e5

```

```

[26]: if os.path.isfile(os.path.join(checkpoint_dir, checkpoint_name)):
    net.load_parameters(os.path.join(checkpoint_dir, checkpoint_name), ctx=ctx)
    print("Parameters loaded")

```

```

[27]: trainer = gluon.Trainer(net.collect_params(), 'adam', {'learning_rate':
    ↳learning_rate})
for e in range(epochs):
    cls_metric = mx.metric.Accuracy()

```

```

    box_metric = mx.metric.MAE()
    train_loss = run_epoch(e, net, train_data, trainer, log_dir,
↪ print_name="train", is_train=True, update_metric=False)
    test_loss = run_epoch(e, net, test_data, trainer, log_dir,
↪ print_name="test", is_train=False, update_metric=True)
    if test_loss < best_test_loss:
        print("Saving network, previous best test loss {:.6f}, current test_
↪ loss {:.6f}".format(best_test_loss, test_loss))
        net.save_parameters(os.path.join(checkpoint_dir, checkpoint_name))
        best_test_loss = test_loss

    if e % print_every_n == 0:
        name1, val1 = cls_metric.get()
        name2, val2 = box_metric.get()
        print("Epoch {0}, train_loss {1:.6f}, test_loss {2:.6f}, test {3}={4:.
↪ 6f}, {5}={6:.6f}".format(e, train_loss, test_loss, name1, val1, name2, val2))

```

[09:56:18] ../src/operator/nm/./cudnn/./cudnn_algoreg-inl.h:97: Running performance tests to find the best convolution algorithm, this can take a while... (set the environment variable MXNET_CUDNN_AUTOTUNE_DEFAULT to 0 to disable)

Saving network, previous best test loss 1000000.000000, current test loss 0.746653

Epoch 0, train_loss 0.858615, test_loss 0.746653, test accuracy=0.949080, mae=0.029872

Saving network, previous best test loss 0.746653, current test loss 0.353118

Saving network, previous best test loss 0.353118, current test loss 0.190814

Saving network, previous best test loss 0.190814, current test loss 0.139893

Saving network, previous best test loss 0.139893, current test loss 0.114793

Saving network, previous best test loss 0.114793, current test loss 0.100580

Epoch 5, train_loss 0.096800, test_loss 0.100580, test accuracy=0.986375, mae=0.022396

Saving network, previous best test loss 0.100580, current test loss 0.091744

Saving network, previous best test loss 0.091744, current test loss 0.084517

Saving network, previous best test loss 0.084517, current test loss 0.080071

Saving network, previous best test loss 0.080071, current test loss 0.075976

Saving network, previous best test loss 0.075976, current test loss 0.072644

Epoch 10, train_loss 0.064722, test_loss 0.072644, test accuracy=0.986552, mae=0.019780

Saving network, previous best test loss 0.072644, current test loss 0.069390

Saving network, previous best test loss 0.069390, current test loss 0.067201

Saving network, previous best test loss 0.067201, current test loss 0.065241

Saving network, previous best test loss 0.065241, current test loss 0.063827

Saving network, previous best test loss 0.063827, current test loss 0.062032

Epoch 15, train_loss 0.053825, test_loss 0.062032, test accuracy=0.987415, mae=0.017978

Saving network, previous best test loss 0.062032, current test loss 0.060555

Saving network, previous best test loss 0.060555, current test loss 0.059710
 Saving network, previous best test loss 0.059710, current test loss 0.058479
 Saving network, previous best test loss 0.058479, current test loss 0.057002

```
[28]: if os.path.isfile(os.path.join(checkpoint_dir, checkpoint_name)):
      net.load_parameters(os.path.join(checkpoint_dir, checkpoint_name), ctx=ctx)
```

Wyniki Funkcja pomocnicza do przewidywania pola ograniczającego

```
[30]: def predict_bounding_boxes(net, image, bb):
      '''
      Given the outputs of the dataset (image and bounding box) and the network,
      the predicted bounding boxes are provided.

      Parameters
      -----
      net: SSD
      The trained SSD network.

      image: np.array
      A grayscale image of the handwriting passages.

      bb: [(x1, y1, x2, y2)]
      A tuple that contains the bounding box.

      Returns
      -----
      predicted_bb: [(x, y, w, h)]
      The predicted bounding boxes.

      actual_bb: [(x, y, w, h)]
      The actual bounding bounding boxes.
      '''
      image, bb = transform(image, bb)

      image = image.as_in_context(ctx[0])
      image = image.expand_dims(axis=0)

      bb = bb.as_in_context(ctx[0])
      bb = bb.expand_dims(axis=0)

      default_anchors, class_predictions, box_predictions = net(image)
      box_target, box_mask, cls_target = net.training_targets(default_anchors,
      ↪class_predictions, bb)
      cls_probs = nd.SoftmaxActivation(nd.transpose(class_predictions, (0, 2,
      ↪1)), mode='channel')
```

```

    predicted_bb = MultiBoxDetection(*[cls_probs, box_predictions,
↪default_anchors], force_suppress=True, clip=False)
    predicted_bb = box_nms(predicted_bb, overlap_thresh=overlap_thres,
↪valid_thresh=min_c, topk=topk)
    predicted_bb = predicted_bb.asnumpy()
    predicted_bb = predicted_bb[0, predicted_bb[0, :, 0] != -1]
    predicted_bb = predicted_bb[:, 2:]
    predicted_bb[:, 2] = predicted_bb[:, 2] - predicted_bb[:, 0]
    predicted_bb[:, 3] = predicted_bb[:, 3] - predicted_bb[:, 1]

    labeled_bb = bb[:, :, 1:].asnumpy()
    labeled_bb[:, :, 2] = labeled_bb[:, :, 2] - labeled_bb[:, :, 0]
    labeled_bb[:, :, 3] = labeled_bb[:, :, 3] - labeled_bb[:, :, 1]
    labeled_bb = labeled_bb[0]
    return predicted_bb, labeled_bb

```

Analiza ilościowa Średnia IOU została obliczona dla zestawu testowego.

```

[31]: def get_iou(box1, box2):
    '''
    Calculate the IOU between two bounding boxes (x, y, w, h)
    '''
    # source: https://www.pyimagesearch.com/2016/11/07/
↪intersection-over-union-iou-for-object-detection/
    x1 = max(box1[0], box2[0])
    y1 = max(box1[1], box2[1])
    x2 = min(box1[2], box2[2])
    y2 = min(box1[3], box2[3])

    inter_area = max(0, x2 - x1 + 1) * max(0, y2 - y1 + 1)

    box1_area = (box1[2] - box1[0] + 1) * (box1[3] - box1[1] + 1)
    box2_area = (box2[2] - box2[0] + 1) * (box2[3] - box2[1] + 1)

    iou = inter_area / float(box1_area + box2_area - inter_area)
    return iou

def calculate_iou(dataset):
    '''
    Iterate through the dataset and calculate the mean IOU between the actual
↪and
    predicted bounding boxes.
    '''
    ious = []
    for i in range(len(dataset)):
        iou_i = 0.0

```

```

count_i = 0
image, bb = dataset[i]
predicted_bb, actual_bb = predict_bounding_boxes(net, image, bb)
# A naive 1-1 bounding box matching algorithm was used. This algorithm
# doesn't account for insertions or deletes and may result in lower IOU
↪ values.
    for predicted_bb_i, actual_bb_i in zip(predicted_bb, actual_bb):
        iou = get_iou(predicted_bb_i, actual_bb_i)
        iou_i += iou
        count_i += 1
    ious.append(iou_i/count_i)
return np.mean(ious)

train_iou = calculate_iou(train_ds)
test_iou = calculate_iou(test_ds)
print("Train iou {} test iou {}".format(train_iou, test_iou))

```

[10:04:15] ../src/operator/nn/./cudnn/./cudnn_algoreg-inl.h:97: Running performance tests to find the best convolution algorithm, this can take a while... (set the environment variable MXNET_CUDNN_AUTOTUNE_DEFAULT to 0 to disable)

Train iou 0.44020686944027226 test iou 0.44010712224485027

Wizualizacja przewidzianych ramek i rzeczywistego tekstu.

```

[32]: figs_to_plot = 2
fig, axs = plt.subplots(figs_to_plot, 2, figsize=(15, 10 * figs_to_plot))
ds = test_ds
for i in range(figs_to_plot):
    n = int(random.random()*len(ds))
    image, bb = ds[n]
    predicted_bb, actual_bb = predict_bounding_boxes(net, image, bb)

    for j in range(actual_bb.shape[0]):
        (x, y, w, h) = actual_bb[j]
        axs[i][0].imshow(image, cmap='Greys_r')
        image_h, image_w = image.shape[-2:]
        (x, y, w, h) = (x * image_w, y * image_h, w * image_w, h * image_h)
        rect = patches.Rectangle((x, y), w, h, fill=False, color="r")
        axs[i][0].add_patch(rect)
        axs[i][0].set_title("BB actual")

    for j in range(predicted_bb.shape[0]):
        axs[i][1].imshow(image, cmap='Greys_r')
        (x, y, w, h) = predicted_bb[j]
        image_h, image_w = image.shape[-2:]
        (x, y, w, h) = (x * image_w, y * image_h, w * image_w, h * image_h)

```

```

rect = patches.Rectangle((x, y), w, h, fill=False, color="r")
axs[i][1].add_patch(rect)
axs[i][1].set_title("BB predicted")

```

